

Albert Query - Agentic Movie Intelligence System

An intelligent conversational agent for querying movie and TV series data using Retrieval Augmented Generation (RAG), multi-tool orchestration, and semantic search.

Python 3.8+

Streamlit 1.51.0

LangChain 1.0.3

LangGraph 1.0.2

ChromaDB 1.3.4

Table of Contents

- [Overview](#)
- [Features](#)
- [Architecture](#)
- [Project Structure](#)
- [Components](#)
- [Workflow](#)
- [Future Improvements](#)
- [Contributors](#)
- [License](#)

Overview

Albert Query is an agentic AI system we developed as part of our M1 project at **Albert School** in collaboration with **Mines Paris - PSL**. The system intelligently answers questions about movies and TV series by orchestrating multiple data sources and tools through a LangGraph-based workflow.

What Makes It Special?

Unlike traditional chatbots, Albert Query:

- **Plans before acting** - Analyzes each question to determine which tools are needed
- **Multi-source intelligence** - Combines SQL databases, vector search, external APIs, and web search
- **Semantic understanding** - Uses OpenAI embeddings to find movies by plot similarity
- **Source attribution** - Always shows where information comes from
- **Context-aware** - Maintains conversation history for follow-up questions

Use Cases

-  **Semantic Search:** "Find me movies about space exploration with AI themes"
-  **Data Analysis:** "How many comedies were released on Netflix after 2020?"
-  **Movie Discovery:** "Show me films similar to Inception"
-  **Trend Analysis:** "What are the top-rated action movies from the 2010s?"

-  **Latest Info:** "What's trending in movies this week?"
-

❖ Features

Core Capabilities

Intelligent Query Planning

- LLM-based planner analyzes questions and conversation history
- Automatically selects optimal tools (SQL, Semantic Search, OMDB, Web)
- Avoids unnecessary API calls for efficiency

Multi-Database SQL Queries

- Comprehensive catalog of 8,000+ movies/shows from Netflix, Disney+, Amazon Prime
- Structured queries with filters (year, genre, rating, type)
- Automatic database schema understanding

Semantic Vector Search

- 114MB of OpenAI embeddings (text-embedding-3-small)
- Find movies by plot descriptions, themes, or similarity
- Natural language queries in English or French

OMDB API Integration

- Enriched movie metadata (actors, awards, ratings, posters)
- IMDb links for detailed information
- Full plot summaries

Web Search

- DuckDuckGo integration for trending topics
- Latest movie news and releases
- Current events in cinema

Source Attribution

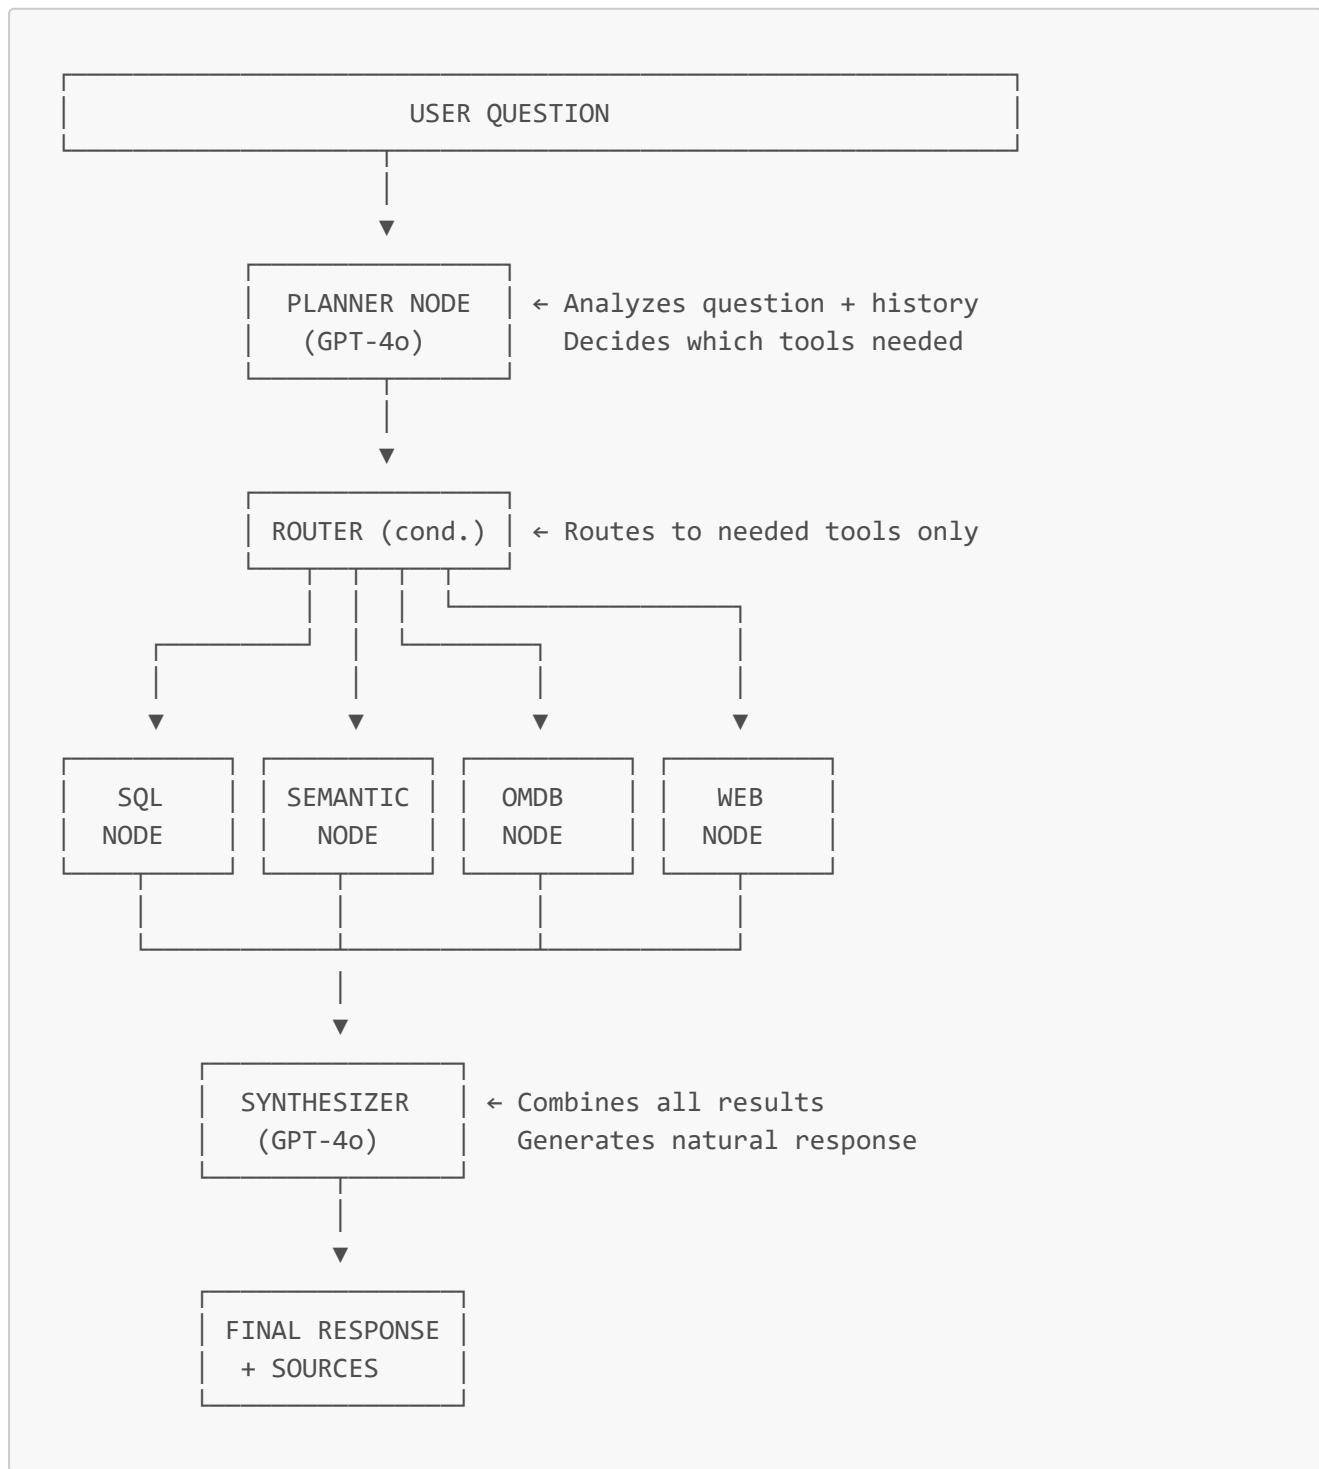
- Tracks all information sources
- Visual display of databases, APIs, and search results used
- Clickable links to IMDb and search results

Conversational Memory

- LangGraph checkpointing with MemorySaver
 - Full conversation history maintained
 - Context resolution ("that movie" → actual movie name)
-

E Architecture

Our system follows an **agentic architecture** using LangGraph to create a stateful, multi-tool workflow:



📁 Project Structure

```
Agentic_Systems_with_RAG_Lamy-Waerniers/
|
|   code/                                     # Main source code (modular
|       architecture)
|           app.py                            # Streamlit UI and main entry
|           point
|           agent.py                         # LangGraph workflow construction
|           nodes.py                         # All workflow nodes (planner,
```

```

SQL, semantic, OMDB, web, synthesizer)
|   ├── tools.py
web search, OMDB API, semantic search)
|   ├── models.py
(AgentState, PlannerOutput, SQLOutput)
|   ├── config.py
keys, paths, LLM instance)
|   ├── utils.py
builder, routing logic)
|   ├── embedding.py
|   └── notebooks/
development
    ├── embedding.ipynb
    ├── SQLdb_creator.ipynb
    └── test_semantic_search.ipynb

    └── data/
        ├── csv_db/
        |   ├── amazon_prime_titles.csv
        |   ├── netflix_titles.csv
        |   └── disney_plus_titles.csv
        ├── databases/
        |   └── movie.db
        ├── vector_database/
        |   ├── chroma.sqlite3
        |   └── 19c0759d-...
        └── memory/
            ├── conversations/
            └── user_profiles/

    └── doc/
        ├── graph_schema.png
        ├── omdb_api_doc.json
        └── OMDB_API_doc.txt

    └── .env
ignored)
└── .gitignore
└── requirements.txt
packages)
└── README.md

# Tool implementations (SQL query,
# Shared type definitions
# Centralized configuration (API
# Helper functions (catalog
# Vector embedding utilities
# Jupyter notebooks for
# Embedding pipeline notebook
# Database creation from CSVs
# Semantic search validation

# Data storage
# Source CSV files
# Amazon Prime catalog (3.9MB)
# Netflix catalog (3.4MB)
# Disney+ catalog (385KB)

# Consolidated SQLite DB (32.9MB)
# ChromaDB persistent storage
# Vector DB metadata (42.7MB)
# Embedding data (114MB)
# Conversation storage

# Documentation
# LangGraph workflow diagram
# OMDB API reference

# Environment configuration (git-
# Git ignore rules
# Python dependencies (223
# This file

```

గ්‍රැෆ් පූර්ව පිළිබඳතාව

Core Modules:

- **config.py** - Central configuration hub
 - API keys (OpenAI, OMDB)
 - Absolute paths to data folders
 - LLM instance (ChatOpenAI)

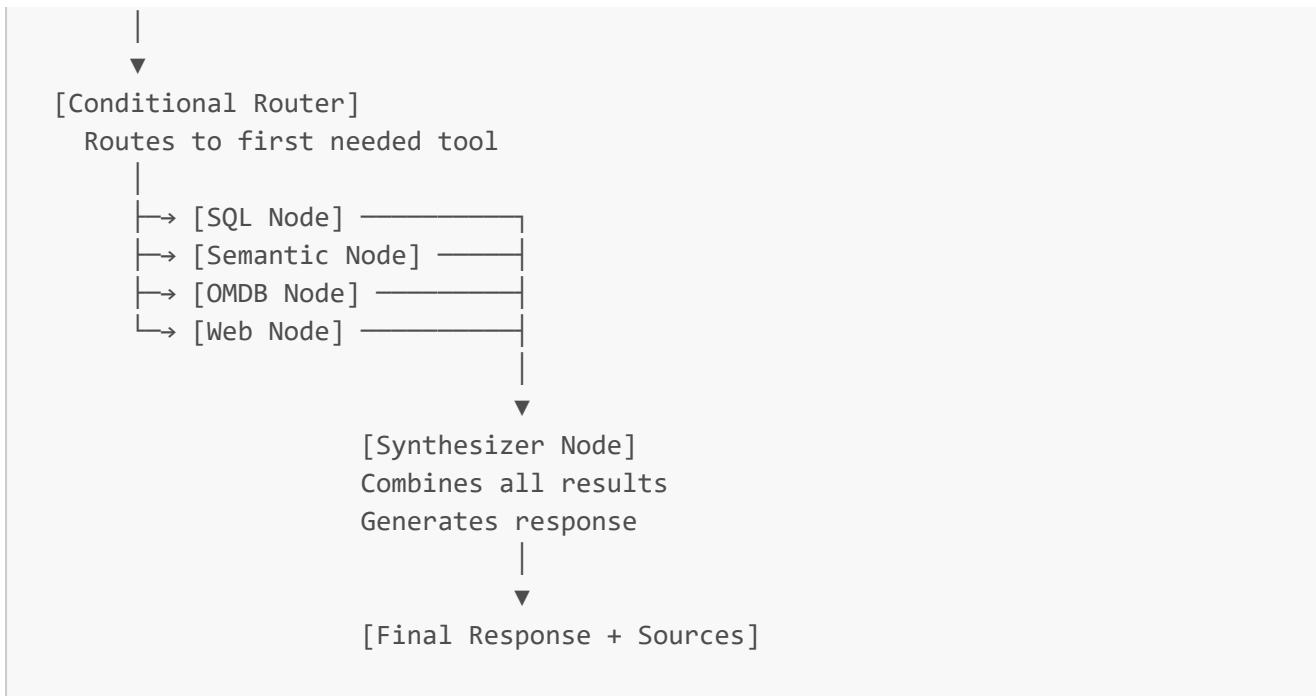
- **models.py** - Shared type definitions
 - `AgentState`: TypedDict defining the workflow state
 - `PlannerOutput`: Pydantic model for planner decisions
 - `SQLOutput`: Pydantic model for SQL execution decisions
 - **tools.py** - Tool implementations
 - `execute_sql_query()`: Query SQLite databases
 - `semantic_search()`: Vector similarity search with ChromaDB
 - `omdb_api()`: Fetch movie metadata from OMDB
 - `web_search()`: DuckDuckGo web search
 - **nodes.py** - LangGraph workflow nodes
 - `planner_node()`: Analyzes question and decides which tools to use
 - `sql_node()`: Generates and executes SQL queries
 - `semantic_search_node()`: Performs vector search
 - `omdb_node()`: Fetches enriched movie data
 - `web_node()`: Searches the web
 - `synthesizer_node()`: Combines results into natural language response
 - **utils.py** - Helper functions
 - `build_db_catalog()`: Introspects database schema
 - `format_catalog_for_llm()`: Formats catalog for LLM prompts
 - Routing functions for conditional edges
 - **agent.py** - LangGraph workflow builder
 - Constructs the StateGraph
 - Defines node connections and routing
 - Compiles workflow with MemorySaver checkpointing
 - **app.py** - Streamlit application
 - UI components (chat interface, source attribution)
 - Session state management
 - Workflow execution and streaming
-

⌚ Workflow

```

User Question
|
↓
[Planner Node]
    Analyzes: question + history
    Decides: which tools needed
    Outputs: needs_sql, needs_semantic, needs_omdb, needs_web

```



🚧 Future Improvements

We've identified several areas for optimization and enhancement. Here's our roadmap:

1. 📁 Catalog Caching System

Problem:

- Database catalog is rebuilt on every app startup
- Slow initialization (~2-5 seconds)
- Redundant SQL queries for schema introspection

Solution:

```

# Implement caching with invalidation detection
def get_or_build_catalog(db_path: str, cache_path: str) -> dict:
    """
    Cache database catalog as JSON
    - Compare file modification times to detect changes
    - Load from cache if DB unchanged
    - Rebuild only when necessary
    """
    
```

Expected Impact:

- ⚡ 10-50x faster startup time
- 🗄️ Reduced SQL queries
- 🎯 Auto-invalidation on schema changes

Implementation:

- Save catalog to `data/databases/catalog_cache.json`
- Include DB file mtime and size for change detection
- Add force-rebuild option for manual invalidation

2. 🧠 Persistent Long-Term Memory

Current State:

- Memory stored in LangGraph's MemorySaver (in-memory only)
- Lost on application restart
- No cross-session learning

Proposed Architecture:

```
# SQLite-based conversation storage
conversations/
    └── user_123/
        ├── session_20250116_001.json      # Conversation history
        ├── session_20250116_002.json
        └── preferences.json               # Learned preferences
    └── user_456/
        └── ...
}

# Conversation schema
{
  "session_id": "20250116_001",
  "user_id": "user_123",
  "timestamp": "2025-01-16T10:30:00Z",
  "messages": [...],
  "topics": ["action movies", "2020s cinema"],
  "preferences_learned": {
    "favorite_genres": ["action", "sci-fi"],
    "preferred_platforms": ["netflix"]
  }
}
```

Features to Add:

- 🗂 Persist conversations to disk (JSON or SQLite)
- 🚙 User-specific history and preferences
- 🔎 Semantic search over past conversations
- 📊 Analytics on user interests
- 💬 Personalized recommendations based on history

Technical Implementation:

- Replace MemorySaver with custom SQLiteCheckpointer
- Add user authentication (see #3)
- Implement conversation summarization for long histories

- Privacy controls (GDPR compliance)

3. User Management & API Key Interface

Current Limitation:

- Single shared API keys in `.env`
 - No multi-user support
 - API costs not attributable to users

Implementation:

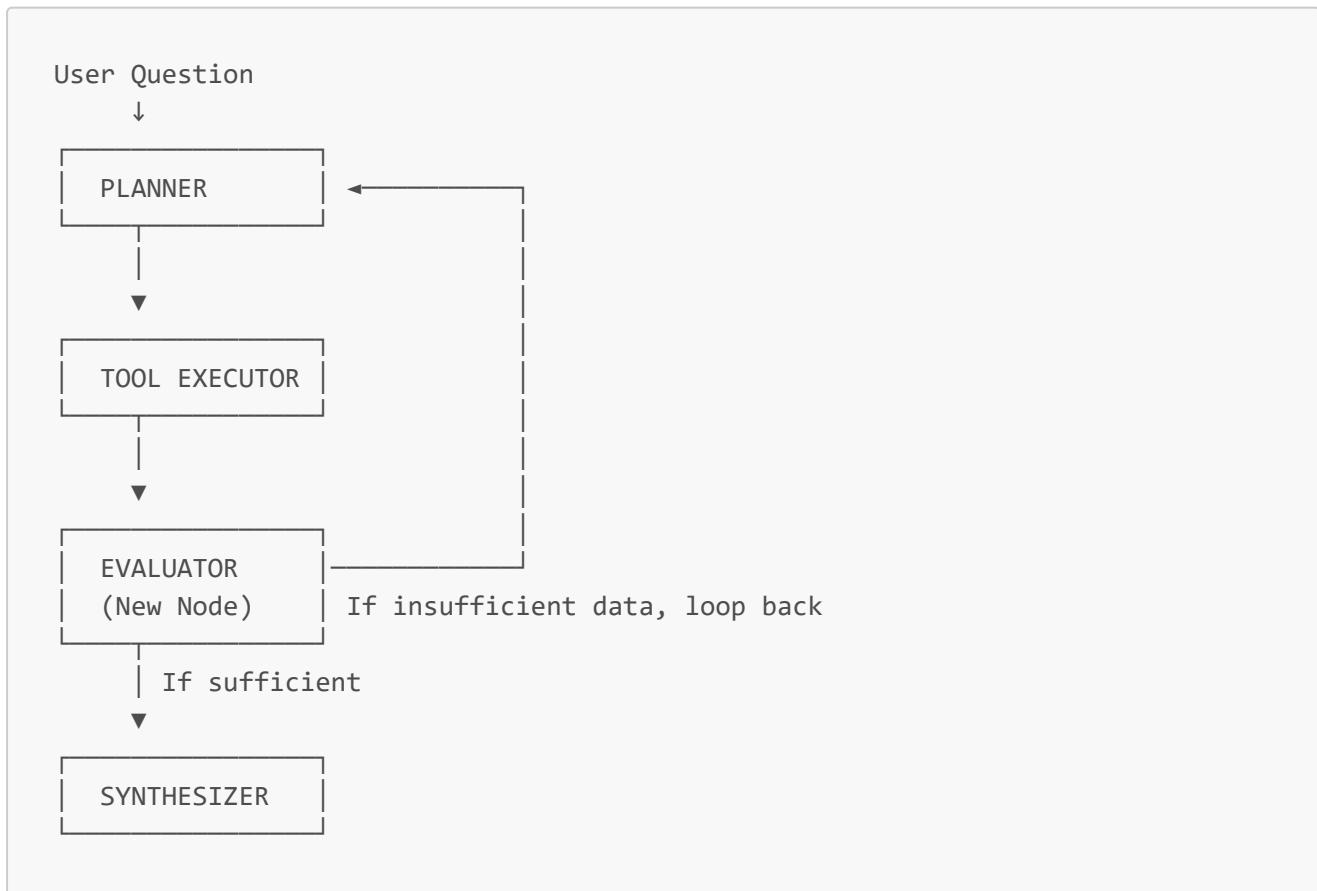
- Streamlit sidebar with settings panel
 - Encrypted key storage per user (keyring library)
 - Session-based authentication
 - Token tracking and cost estimation
 - Rate limiting per user

4. Workflow Enhancement: Planner Loop

Current Issue:

- Linear workflow: Planner → Tools → Synthesizer → End
 - No feedback loop if initial plan was insufficient
 - Cannot self-correct or ask for more tools

Proposed Architecture:



5. Embedding Quality Improvements

Current Limitations:

- Similarity scores often low (<50%)
- Movie descriptions are single sentences only
- No chunking strategy
- Basic embedding model (text-embedding-3-small)

Improvement Strategies:

A. Enhance Movie Description Quality (Priority #1)

Currently, we only embed the plot description field from databases (single sentence).

Use APIs to enrich our database with way longer movie descriptions.

Expected Impact:

-  More contextual embeddings
-  Better similarity scores (10-20% improvement)
-  Genre/cast matching in semantic search

6. Structured Output Enforcement

Problem:

- JSON parsing errors possible in synthesizer
- Inconsistent response formats
- Difficult to extract structured data

7. Token Optimization

Current Issues:

- Prompts are verbose (500-800 tokens each)
- Full database catalog sent to planner (1000+ tokens)
- Conversation history grows unbounded

Optimization Strategies:

A. Prompt Compression

B. Catalog Summarization

C. Conversation Summarization

D. Lazy Loading

8. UI/UX Enhancements

Proposed Features:

- **Results Table View** - Toggle between chat and table display
- **Movie Cards** - Rich display with posters, ratings, cast
- **Query Statistics** - Show token usage, cost, response time
- **Dark Mode** - Theme switching
- **Export Results** - Download conversations as JSON/CSV
- **Voice Input** - Speech-to-text for queries
- **Multi-language Support** - Full i18n for French/English
- **Mobile Responsiveness** - Optimize for mobile devices
- **Keyboard Shortcuts** - Power user features

9. Testing & Quality Assurance

Current Gap: No automated tests

Coverage Goals:

- Unit tests for each node
- Integration tests for workflow
- Performance benchmarks
- Regression tests for common queries

10. Security & Privacy

Enhancements Needed:

- API key encryption at rest
- Input sanitization (SQL injection prevention)
- PII detection and redaction in conversations
- Audit logging for all queries
- Rate limiting and abuse prevention
- HTTPS enforcement in production
- Content filtering for inappropriate queries

11. Performance & Scalability

Optimization Opportunities:

- **Async Tool Execution** - Run SQL, Semantic, OMDB in parallel
- **Result Caching** - Cache common queries (Redis)
- **Vector Index Optimization** - Use HNSW parameters tuning
- **Database Indexing** - Add indexes on common query columns
- **Connection Pooling** - Reuse DB connections
- **Deployment** - Docker + cloud hosting (AWS/GCP)
- **CDN Integration** - Cache static assets

12. Analytics & Monitoring

Tracking Metrics:

- Query latency by tool type
- Cost per query (token usage)
- Tool selection accuracy (planner effectiveness)
- User engagement metrics
- Error rates and types
- Most common queries and topics
- Semantic search quality metrics

Contributors

This project was developed as part of our Master's degree at **Albert School X Mines Paris - PSL**.

Team:

- Vincent Lamy & Alexandre Waerniers

Institution:

- Albert School (Paris, France)
 - Mines Paris - PSL
-

License

This project is licensed under the MIT License - see the LICENSE file for details.

★ If you found this project useful, please consider giving it a star! ★